



Aalto University  
School of Science

# CS-E4110

## Concurrent Programming

Week 3 – Exercise session

2021-11-19

14:15 – 16:00

Jaakko Harjuhahto

# Today

- Semaphores
- Standard APIs
- Reactor assignment
- General Q&A

# Semaphores

# Semaphore Semantics

- A semaphore is a compound type of a non-negative integer and a set of waiting processes
  - The value of the integer is often called the number of permits or the number of tickets
- The two operations `wait` and `signal` are atomic and mutually exclusive
  - With `P(s)` a process decrements  $s$  ( $> 0$ ) by one. If  $s$  is 0, the process is put to wait “in `P(s)`” until  $s$  is positive so that  $s$  can be decremented and the process may proceed
  - With a `V(s)` a process increments  $s$  by one. If  $s$  was 0, and there are one or more processes waiting “in `P(s)`”, one of them may complete `P(s)` and proceed

# Terminology

Decrement	Increment	
<code>wait(s)</code>	<code>signal(s)</code>	Standard synchronization terminology
<code>P(s)</code>	<code>V(s)</code>	Classic shorthand from Dijkstra
Prolaag	Verhoog	E.W.Dijkstra's original terminology
<code>s.acquire()</code>	<code>s.release()</code>	Java API terminology
<code>down(s)</code>	<code>up(s)</code>	Linux kernel

Even more variants exist...

# Example: Simple Critical Section

- Since the semaphore only has a single permit available, if a thread is inside the critical section, other threads encountering the `wait()` will be placed in the wait set
  - `signal()` will wake a thread, which can then acquire a permit and proceed

```
val sem = new Semaphore(1)

sem.wait()    //permits -1
//critical section
sem.signal()  //permits +1
```

# Example: Semaphore With JVM Monitors

```
class MySemaphore(var tickets: Int) {  
  
    def wait(): Unit = ???  
  
    def signal(): Unit = ???  
  
}
```

# Gotchas

- The semaphore can be an object with methods
  - E.g. `semInstance.wait()`
- The semaphore can be a struct that is operated on by static methods
  - E.g. `sem_wait(semInstance)`
- The initial ticket count of a general semaphore is not an upper limit for tickets or a maximum capacity
  - Example: It is valid to initialize a semaphore with a ticket count of 2 and then increase it to 5 by making three calls to `signal()`



# Standard APIs

# Scala & Java Standard APIs

- Most of the standard library is **not** thread safe
  - Default assumption: calling API methods from multiple threads will break something, especially if calls happen concurrently
    - Lack of mutual exclusion
    - Lack of happens-before
  - The few thread safe exceptions are documented
- We can always implement a thread safe variant of an API class by writing a new class that wraps the methods of the original class in synchronized methods

# Standard Library

- On this course, we **implement** concurrent code using threads, JVM monitors and other language level primitives to **learn** concurrent programming
- For practical day-to-day usage outside this course, the standard library offers APIs for concurrency
  - `java.util.concurrent`
  - `scala.concurrent`
- The same logic applies to other languages: there is a good chance that a standard library or a common framework might already have what you need

# Examples

- Concurrent collections (lists, queues, maps, sets etc.)
- Atomic variables
  - e.g. `AtomicInteger`, `AtomicReference`
- Explicit locks
  - e.g. `ReadWriteLock`
- Executors
  - Don't explicitly create threads, instead submit items of work, e.g. tasks, to a service that will take care of executing these items
- Synchronization constructs
  - `CyclicBarrier`, `CountDownLatch`, `Phaser`
- We will discuss parallel collections, futures and promises in more detail later on during this course

# Reactor Assignment



# Reactor Assignment

- Course mini-project in three phases
  - Objective: implement a small networked game server
  - Problem: network programming is inherently concurrent, as we must listen to multiple clients concurrently
  - Solutions: use a design pattern to separate concurrency issues from the game logic
- A step up in complexity compared to A+ exercises
- Tasks define an API and behaviour
  - You may choose how to implement your own solution
  - Several good solution archtypes for the problems exist

# The Reactor Pattern

- Design patterns are re-usable solutions to common problems in software design
- The reactor is a mechanism for changing an inherently concurrent problem into an asynchronous but sequential problem
  - All the necessary threading and locking is implemented inside the reactor and the application built on top of the reactor can be written as sequentially executed event handlers
  - Our use case: responding to multiple network clients
- Similar designs: JavaScript interpreter event loop, UI toolkit dispatchers, etc...

# Common Terminology

- Handles represent sources of data
  - Reading the handle will produce an element of data
  - May block while reading, until data becomes available
    - E.g. a network client sends data
- Events contain an element of data and a handler
- Handlers are code which accepts the element of data and does something with it, typically changes program state
  - E.g. "when a new player connects (= the event), register the player with the game, so that they will get updates from the game server"



# Task A – BlockingEventQueue

- Theme: JVM monitors and synchronization
- Implement a FIFO bounded buffer for storing objects
  - Blocks if trying to get from an empty buffer
  - Blocks if trying to put into a full buffer
  - The buffer contains events of an generic type A
    - I.e. class `BlockingEventQueue[A]`
- The base API defines four methods:
  - `enqueue()`, `dequeue()`, `getSize()` and `getCapacity()`
    - Only `enqueue()` and `dequeue()` are necessary for Task-B
- Two optional features give extra points
  - `getAll()` method and notification efficient design

# Task-A Degrees of Difficulty

Features	Points	Difficulty
Base API	60	★
API + getAll	60 + 15	★★
API + notification efficient	60 + 15	★★
API + getAll + notification efficient	60 + 15 + 15	★★★

Beware: Implementing the optional goals can break the base API implementation and introduce new bugs

# Task B – Dispatcher

- Theme: thread management and design patterns
- The Dispatcher maintains a registry of handles
- It is responsible for starting threads to read handles when they are registered
  - Also, must ensure that the started threads terminate correctly
- Two fundamental requirements:
  - The dispatcher must be able to read all the registered handlers concurrently
  - The event loop must process events from handles sequentially, one at a time, in a single thread

# Task C – The Game

- Theme: asynchronous/callback programming
- Implement the Hangman game on top of the reactor
  - Wikipedia: [https://en.wikipedia.org/wiki/Hangman\\_\(game\)](https://en.wikipedia.org/wiki/Hangman_(game))
  - The ruleset is implemented for you in the template
- You must implement game logic by defining handlers that change game state when certain events occur
  - Events are, for example: player joins the game, player makes a guess, etc
  - As the game is a network server, you don't know the order incoming of events

# The Template

- A ZIP package with Scala classes you must implement
  - Includes the necessary API classes
  - Utility classes you can freely use as a part of your Task-C
- The template contains minimal unit tests
  - Only check that basic semantics work, the tests do not exhaustively test concurrency behavior
  - You can of course implement more tests...
- The same instructions as in A+

# Assignment Practicalities



# Submission

- Each task is returned to A+ as a single `.scala` file
  - If your solution needs multiple classes, implement them in the same source file
- Distinct deadline for each ask
  - Task-A Monday 29th of November at 14:00
  - Task-B Monday 6th of December at 14:00
  - Task-C Monday 13th of December at 14:00
- Only the last submission at the time of the DL is graded
  - You can make up to 100 submissions before the DL

# Grading (1/2)

- The assignment submissions are graded manually
  - The target for publishing the grading is the next Friday
- After the grading, you receive feedback and points
  - The feedback gives you a breakdown of how the point score is composed and refers to specific issues with the submission
- Each task is worth 100 points
  - Up to 90 points for the implementation code
  - Up to 10 points for documentation (= comments)



# Grading (2/2)

- For each part, the submission starts with a base level of points
  - Base is 60 points for Task-A and 90 points for Tasks B and C
  - For Task-A, optional features give +15 and +15 points
- Errors, bugs, design flaws and deviations from the specification **deduct** points, depending on the severity
  - For example, if the design can potentially deadlock: -25 points
  - Beware stacking bugs

# Documentation

- Part of the grading criteria
  - The comments should explain why your implementation is correctly synchronized as defined in the Java Language Specification
  - Help the TAs understand your solution
  - In borderline cases, the documentation is used to determine if a concurrency problem was the result of a coding error (minor point penalty) or a design flaw (larger penalty)
- Documentation gives up to 10 points for each Task

# Q&A