



Aalto University
School of Science

Concurrency frameworks

CS-E4110 Concurrent Programming
Autumn 2021 (II), week 4

Vesa Hirvisalo

Week 4 in a nutshell

- Computing in a large scale
 - Scalability, resilience, efficiency
 - Distributed systems, datacenters, etc.
- Introduction to concurrency frameworks
 - Working on a higher abstraction level
 - Note: usually does not exclude using the lower levels!
- Akka actors
 - A message passing system for Scala
 - The system supports actor-local shared memory computing
 - With (some) fault tolerance and natural distribution

Computing in large scale

Distributed systems

- In larger scale we need
 - Scalability, resilience, efficiency, etc.
=> Several computers / computing nodes
Highly concurrent distributed systems
- Network environments and distributed environments
 - Network: the individual computers a visible to the user
 - Distributed: viewed as a single environment by the user
- Distributed computing is challenging
 - There is no global state (by definition)
 - The system state is update by exchanging message
 - Without a global state, understanding of correctness is hard
(we will omit a detailed view into these problems)

Distribution and synchronization

- Synchronization is about the (correct) execution order(s)
 - We allow only those with the correct result
 - The execution order is related to time
- What is the time now?
 - The answer comes always with a delay
 - Distributing a clock is difficult
 - And similarly, being synchronous is difficult
- The correct order of events?
 - Assuming no global time (in a system)
 - The system cannot tell (with means of the system itself)
 - Only partial orders are possible

Datacenters as computers

- Efficient and scalable computing
 - Need plenty of resources => make computing a utility
 - Energy consumption => highly flexible and concurrent code
- Typical datacenters are warehouse size computers
 - Built around a set of networks
- Usually, there is
 - A set of networks
 - A hierarchy of memory and storage
 - The processing units
 - In addition
 - Plenty issues on power, cooling, etc.
 - It really is a house!

Introduction to frameworks

Shared memory frameworks

- Traditionally uses locks (etc.) to synchronize and protect data structures
 - If a thread crashes while holding a lock, other threads will block on the lock
 - Concurrency is limited to a single machine shared memory
- A number of frameworks for parallelization
 - More abstract than raw threads, e.g., OpenMP
- Distribution, fault-tolerance, etc.?
 - No single (or direct) way to do such
 - E.g., partially solving fault-tolerance: use short lock regions that are robust against exceptions, release held locks in exception handlers

Message passing frameworks

- Have an abstract and integrated approach
 - Instead of simply having only some send/receive methods
 - E.g., MPI used for high performance computing
- Message passing frameworks are a natural choice for distributed computing
 - Distributed systems are often defined as systems that base their communication on messages (shared channels instead of shared memory)
- However, some things are not easy
 - Especially, understanding system state
 - E.g., detecting when an algorithm has terminated

Functional programming frameworks

- Several aspects can be implemented by a framework
 - Dataflow, synchronization, fault tolerance
 - E.g., recomputing a function on unmodified input data will yield the same result
 - Natural way to parallelization and load balancing
 - Some frameworks: MapReduce, Apache Spark, Scala parallel collections
- However, functions are by default computed in isolation
 - Communication heavy algorithms such as many iterative algorithms can become slow
 - Lack of data mutations can make some algorithms slower

Actor frameworks

- Actor frameworks are a mix between shared memory and message passing concurrency
 - Each actor has its own mutable state
 - Actors communicate by sending immutable messages
 - The actor frameworks give ways to do implement application specific load balancing and fault tolerance strategies with help from the framework
- Actors are made as lightweight as possible
 - Much more lightweight than threads
 - The message send/receive is location transparent
 - The same messaging primitives are used for both local messages and messages over the network

Reactive manifesto

- The Manifesto (www.reactivemanifesto.org) says that reactive systems are:
 - Responsive: Systems respond in a timely manner
 - Resilient: The system stays responsive in the face of failure
 - Elastic: The system stays responsive under varying workload
 - Message driven: Systems rely on asynchronous message passing: loose coupling, isolation, location transparency, delegates errors as messages

Akka actors (in Scala)

The Akka Actor framework

- Design heavily influence by telecommunications software design style - Erlang
 - Used to implement reactive server software that responds to client requests
- Very light weight actors
 - Usually one actor used for each incoming request / connection
 - The Akka documentation is available from [akka.io](https://doc.akka.io/docs/akka/current/index-classic.html):
<https://doc.akka.io/docs/akka/current/index-classic.html>
 - Note that there are several versions of Akka
- Actor framework for Scala and Java
 - **Actors are not objects**
 - Actors have no public methods or data
 - The only way using them is by sending messages

An example

Defining a simple Actor (Classic Akka)

```
1 import akka.actor.Actor
2 import akka.actor.Props
3 import akka.event.Logging
4
5 class MyActor extends Actor {
6   val log = Logging(context.system, this)
7
8   def receive = {
9     case "test" => log.info("received_test")
10    case _ => log.info("received_unknown_message")
11  }
12 }
```

- Note that `this` and `self` are different
 - `self` is used to refer to the `ActorRef` of the Actor
 - A received message has a sender actor reference attached

Akka Futures*

- In futures are a generic concurrency primitive
 - They are for waiting thing that will eventually become available
- In Akka: asynchronous waiting for message reply
- A timeout is usually used to limit the wait
 - Otherwise you should be sure
 - `Await.result` obtains the result of a future
 - `Await.ready` can be used to check is the result ready
- Notice that blocking will prohibit other actors from running on the same thread
 - I.e., blocking is not recommended for performance reasons

**) Next week, we will discuss futures in a more generic way*

Sending messages

- `tell` sends a message asynchronously and returns immediately (basically: `targetActorRef ! Message`)
- `ask` is a way to implement sending a message and generating a `Future` representing a possible reply
 - *E.g.:* `val future = myActor ? "hello"`
 - A timeout for the reply can be specified with the `ask` method
 - `val f = myActor.ask("msg")(5 seconds)`
- At the receiver end for each received message, the `receive` is invoked

Processing messages

- Actors are event driven: Activated at the receive time
 - Multiple messages are queued per actor (FIFO Mailbox)
 - Actors should be using non-blocking I/O. If blocking I/O is needed, special care must be taken not to block other actors from proceeding during the I/O wait
 - Actors themselves are single threaded by design - If concurrency is needed, deploy several actors
- Fault-tolerance through the “Let it crash”-philosophy
 - Isolate the failure and continue with the non-faulty parts, restarting failed subsystems when needed
- Actors operate on local state
 - not to be exposed outside, no internal locking, etc.

Children and supervision

- An Actor can create children
 - The spawned actors form a hierarchy, where each actor is supervising its child actors
 - If errors happen in a child actor, its parent is notified, and it can take supervision action (e.g., restarting the child actor)
- Fault-tolerance is implemented by supervision strategies
 - Each actor has a strategy for supervising its children
 - If different strategies are needed for different children, one needs to create an additional layer of actors to do the different supervision strategies for different children
- As a whole, a system is defined
 - This includes multiple aspects (e.g., scheduling)

Actor Systems

An Example of creating a system (Classic Akka)

```
1  // Modified from "Scala Cookbook":
2  import akka.actor.Actor
3  import akka.actor.ActorSystem
4  import akka.actor.Props
5  class HelloActor extends Actor { // Code removed
6  }
7
8  object Main extends App {
9      // an actor needs an ActorSystem
10     val system = ActorSystem("HelloSystem")
11     // create and start the actor
12     val helloActor = system.actorOf(Props[HelloActor], name = "helloactor")
13     // shut down the system
14     system.shutdown
15 }
```

Supervision strategies

- Main supervision strategies for different errors (e.g., thrown exceptions) are:
 - Resume child, keeping accumulated internal state
 - Restart child, clearing accumulated internal state
 - Stop the child permanently
 - Escalate the failure, failing the actor itself
- Note the hierarchical nature of actors
 - E.g., stopping an actor will all stop all its children
- Some notes
 - Supervision messages have their own system mailboxes and are not ordered with the normal messages
 - Messages `Kill`, `stop` and `PoisonPill` do differ

Location transparency

- Actor references
 - Can refer to actors in remote computers
 - Each actor has an actor path consisting of the hierarchical naming of actor instances
- Everything is distributed by default
 - The distribution is transparent to the programmer
 - All message passing is asynchronous
 - Because messages pass over real network connections, the possibility of losing messages is much higher than within a single machine
 - Sometimes message passing can be optimized by the framework if the sender and receiver are on the same computer

Akka and Java Memory Model

- The happens-before guarantees made by Akka:
 - The send of a message by an actor happens-before the receive of that message by the same actor
 - Processing of one message happens-before processing of the next message by the same actor
 - Thus there is no need to mark internal state volatile
 - However, the messages sent to other actors need to be immutable and properly constructed to not cause problems in data transfer between actors
- Properly constructed immutable data is safe to share between threads

Akka Message Delivery

- Akka takes care of the following:
 - At-most-once delivery
 - Never deliver twice but messages **can get lost!**
 - Applications have to take care of the lost ones
 - There are libraries to help implementing at-least-once delivery
 - Best effort message delivery (i.e., go for low overhead)
 - Message ordering by send-receive pair
- Akka persistence
 - A collection of libraries allowing actors to persist internal state in an external database to allow for better fault tolerance
 - There are also many other tools available

Related readings and exercises

=> Check A+